

# DMC2016

April 22, 2016

```
In [1]: import graphlab
```

A newer version of GraphLab Create (v1.8.5) is available! Your current version is v1.8.4.

You can use pip to upgrade the graphlab-create package. For more information see <https://dato.com/produ>

```
In [2]: url = 'C:/Users/yehplus.bl0076-xx/Dropbox/DMC2016/featureMatrix_v01/train.csv'
```

```
In [2]: train_url_mac = '/Users/yaosuhan/Desktop/DMC_2016_local/featureMatrix_v03/train.csv'
        test_url_mac = '/Users/yaosuhan/Desktop/DMC_2016_local/featureMatrix_v03/test.csv'
```

```
In [3]: raw = graphlab.SFrame.read_csv(train_url_mac)
```

2016-04-22 16:21:11,138 [INFO] graphlab.cython.cy\_server, 176: GraphLab Create v1.8.4 started. Logging:

Finished parsing file /Users/yaosuhan/Desktop/DMC\_2016\_local/featureMatrix\_v03/train.csv

Parsing completed. Parsed 100 lines in 0.940106 secs.

This non-commercial license of GraphLab Create is assigned to yeh@iastate.edu and will expire on March 01, 2017.

Read 42821 lines. Lines per second: 23639

Finished parsing file /Users/yaosuhan/Desktop/DMC\_2016\_local/featureMatrix\_v03/train.csv

Parsing completed. Parsed 315445 lines in 6.87474 secs.

Inferred types from first line of file as

column\_type\_hints=[int,int,float,float,float,float,float,float,float,float,float,float,float,float,float]

If parsing fails due to incorrect types, you can correct

the inferred type list above and pass it to read\_csv in

the column\_type\_hints argument

```
In [4]: test = graphlab.SFrame.read_csv(test_url_mac)
```

Finished parsing file /Users/yaosuhan/Desktop/DMC\_2016\_local/featureMatrix\_v03/test.csv

Parsing completed. Parsed 100 lines in 0.920094 secs.

Read 42272 lines. Lines per second: 23821.7

Finished parsing file /Users/yaosuhan/Desktop/DMC\_2016\_local/featureMatrix\_v03/test.csv

Parsing completed. Parsed 117779 lines in 2.83133 secs.

```
-----
Inferred types from first line of file as
column_type_hints=[int,int,float,float,float,float,float,float,float,float,float,float,float,float,float,float]
If parsing fails due to incorrect types, you can correct
the inferred type list above and pass it to read_csv in
the column_type_hints argument
-----
```

```
In [5]: train,valid = raw.random_split(.8)
```

```
In [6]: train_new = train.dropna('rrp')
```

```
In [7]: model_boosted_tree = graphlab.boosted_trees_classifier.create(raw,target='returnQuantity',
                                                                    max_iterations=100)
```

Boosted trees classifier:

```
-----
Number of examples      : 299782
```

```
Number of classes      : 6
```

```
Number of feature columns : 98
```

```
Number of unpacked features : 98
```

```
+-----+-----+-----+-----+-----+-----+
| Iteration | Elapsed Time | Training-accuracy | Training-log_loss | Validation-accuracy | Validation-loss |
+-----+-----+-----+-----+-----+-----+
| 1         | 6.143165     | 0.677282          | 1.384584          | 0.674137           | 1.385116        |
| 2         | 12.300334    | 0.685151          | 1.165891          | 0.680585           | 1.167275        |
| 3         | 17.550624    | 0.687730          | 1.022610          | 0.683139           | 1.024590        |
```

4	23.124646	0.688781	0.922301	0.684352	0.924375
5	28.547202	0.689835	0.849165	0.685373	0.851561
6	34.649265	0.690942	0.794231	0.685692	0.796904
7	39.902395	0.691563	0.752350	0.686522	0.755142
8	45.074196	0.692076	0.719843	0.687161	0.722942
9	50.376700	0.692810	0.694169	0.687672	0.697372
10	56.105041	0.693150	0.673672	0.688757	0.677155
11	61.378295	0.693431	0.657115	0.688885	0.660906
12	66.792222	0.693794	0.643464	0.688246	0.647353
13	72.170216	0.694078	0.632287	0.688821	0.636365
14	78.533600	0.694731	0.622610	0.689651	0.626923
15	83.760975	0.695375	0.614479	0.689842	0.619016
16	88.962506	0.696216	0.607440	0.690417	0.612586
17	94.826288	0.697043	0.601300	0.691183	0.606903
18	99.933049	0.697554	0.596267	0.691119	0.602256
19	105.055165	0.698124	0.592009	0.691502	0.598307
20	110.149930	0.698748	0.588104	0.691758	0.594956
21	115.314138	0.699412	0.584838	0.692013	0.592194
22	120.700775	0.699762	0.582053	0.691119	0.589695
23	125.932835	0.700352	0.579625	0.691375	0.587644

24	130.981896	0.701169	0.577279	0.691630	0.585890
25	136.029217	0.701546	0.575090	0.692013	0.584350
26	142.437439	0.702010	0.573405	0.692268	0.582991
27	149.884867	0.702594	0.571730	0.692013	0.581810
28	154.870633	0.702774	0.570557	0.691694	0.580860
29	159.988958	0.703084	0.569491	0.692077	0.580004
30	165.074487	0.703378	0.568426	0.691821	0.579440
31	170.093345	0.703781	0.567249	0.691949	0.578595
32	175.069587	0.704092	0.566467	0.691949	0.578025
33	180.185950	0.704639	0.565464	0.691566	0.577816
34	185.447219	0.704762	0.564938	0.691949	0.577424
35	190.470163	0.704979	0.564351	0.691438	0.577200
36	195.855027	0.705249	0.563711	0.692141	0.576851
37	202.877823	0.705466	0.562991	0.691949	0.576691
38	208.173630	0.705980	0.562320	0.691438	0.576421
39	213.812139	0.706280	0.561738	0.691502	0.576268
40	218.980700	0.707054	0.560921	0.691885	0.575813
41	224.117475	0.707471	0.560173	0.691949	0.575667
42	229.222973	0.707788	0.559714	0.692332	0.575576
43	234.045439	0.707904	0.559497	0.692268	0.575544

44	239.587690	0.708091	0.559095	0.692460	0.575410
45	245.194276	0.708762	0.558444	0.692588	0.575341
46	250.040814	0.708862	0.558201	0.692460	0.575266
47	255.099921	0.709275	0.557560	0.692141	0.575034
48	260.092780	0.709766	0.556994	0.692396	0.574868
49	265.858663	0.709946	0.556520	0.692588	0.574680
50	270.869963	0.710153	0.556127	0.692843	0.574599
51	275.941008	0.710500	0.555625	0.692715	0.574445
52	280.868849	0.710873	0.555153	0.691949	0.574229
53	285.871816	0.711017	0.554923	0.692652	0.574188
54	290.835253	0.711153	0.554589	0.692779	0.574247
55	296.425253	0.711207	0.554390	0.692588	0.574250
56	301.962023	0.711607	0.553888	0.692779	0.574194
57	306.911252	0.711817	0.553618	0.692652	0.574194
58	311.734121	0.712121	0.553229	0.692779	0.574068
59	316.525557	0.712484	0.552803	0.692396	0.573865
60	321.565492	0.712605	0.552493	0.692652	0.573926
61	326.397340	0.712701	0.552296	0.692843	0.573855
62	331.896544	0.712735	0.552188	0.692779	0.573867
63	337.672233	0.712938	0.551965	0.692843	0.573877

64	344.054399	0.713245	0.551494	0.693290	0.573865
65	350.544139	0.713335	0.551367	0.693162	0.573879
66	355.470647	0.713575	0.550948	0.692971	0.573724
67	360.253388	0.713619	0.550777	0.693162	0.573739
68	365.092991	0.713999	0.550291	0.692588	0.573850
69	370.457278	0.714306	0.549817	0.692141	0.574001
70	376.104018	0.714549	0.549359	0.692971	0.574129
71	380.907227	0.714809	0.549039	0.692907	0.574196
72	385.968860	0.715106	0.548775	0.692779	0.574342
73	391.108487	0.715523	0.548208	0.693162	0.574182
74	396.664278	0.715920	0.547652	0.693609	0.573976
75	405.421027	0.716501	0.547125	0.694120	0.573807
76	410.157882	0.716714	0.546826	0.694439	0.573727
77	415.260332	0.717114	0.546393	0.693609	0.573691
78	420.387915	0.717955	0.545794	0.694120	0.573640
79	425.337819	0.718112	0.545469	0.693673	0.573725
80	430.251949	0.718122	0.545393	0.693865	0.573754
81	435.414487	0.718375	0.544974	0.693801	0.573795
82	440.471961	0.718782	0.544535	0.693481	0.573799
83	445.623856	0.719039	0.544100	0.693226	0.573696

84	450.701118	0.719396	0.543548	0.692907	0.573571
85	456.314214	0.719716	0.543196	0.693545	0.573674
86	461.257251	0.719850	0.542991	0.693737	0.573671
87	466.218663	0.719923	0.542844	0.693801	0.573685
88	470.929789	0.720270	0.542394	0.694375	0.573541
89	475.879544	0.720343	0.542196	0.694184	0.573580
90	480.925216	0.720614	0.541870	0.695141	0.573482
91	487.447120	0.720881	0.541561	0.695014	0.573524
92	494.186639	0.720904	0.541435	0.695205	0.573553
93	499.130905	0.721241	0.540946	0.695014	0.573498
94	504.665177	0.721534	0.540485	0.694631	0.573511
95	509.164459	0.721721	0.540150	0.695014	0.573668
96	514.038849	0.721915	0.539905	0.694886	0.573783
97	519.635206	0.722335	0.539512	0.694503	0.573916
98	526.408947	0.722635	0.539059	0.694248	0.574023
99	531.175425	0.722992	0.538681	0.694631	0.574068
100	536.971533	0.723296	0.538177	0.695014	0.573845

+-----+-----+-----+-----+-----+-----+

PROGRESS: Creating a validation set from 5 percent of training data. This may take a while.  
 You can set `validation_set=None` to disable validation tracking.

In [8]: `model_boosted_tree.show()`

Canvas is accessible via web browser at the URL: <http://localhost:52029/index.html>  
Opening Canvas in default web browser.

```
In [9]: predictions = model_boosted_tree.predict(test)
```

```
In [13]: model_boosted_tree.evaluate(test)
```

```
Out[13]: {'accuracy': 0.6952937280839538,  
          'auc': 0.9208015979757028,  
          'confusion_matrix': Columns:  
              target_label      int  
              predicted_label    int  
              count              int
```

Rows: 17

Data:

target_label	predicted_label	count
3	2	2
5	1	3
3	4	1
2	1	86
0	2	12
2	2	32
4	0	1
1	0	15618
3	0	6
3	1	3

[17 rows x 3 columns]

Note: Only the head of the SFrame is printed.

You can use `print_rows(num_rows=m, num_columns=n)` to print more rows and columns.,

```
'f1_score': 0.27553175013437753,  
'log_loss': 0.5787037434119917,  
'precision': 0.35746711339564347,  
'recall': 0.2646634589302001,  
'roc_curve': Columns:
```

```
    threshold      float  
    fpr          float  
    tpr          float  
    p            int  
    n            int  
    class        int
```

Rows: 600006

Data:

threshold	fpr	tpr	p	n	class
0.0	1.0	1.0	56978	60801	0
1e-05	1.0	1.0	56978	60801	0
2e-05	1.0	1.0	56978	60801	0



	3e-05		1.0		1.0		56978		60801		0	
	4e-05		1.0		1.0		56978		60801		0	
	5e-05		1.0		1.0		56978		60801		0	
	6e-05		1.0		1.0		56978		60801		0	
	7e-05		1.0		1.0		56978		60801		0	
	8e-05		1.0		1.0		56978		60801		0	
	9e-05		1.0		1.0		56978		60801		0	

+-----+-----+-----+-----+-----+-----+

[600006 rows x 6 columns]

Note: Only the head of the SFrame is printed.

You can use `print_rows(num_rows=m, num_columns=n)` to print more rows and columns.}

In [15]: `model_boosted_tree.summary()`

Class : BoostedTreesClassifier

Schema

-----

Number of examples : 299782

Number of feature columns : 98

Number of unpacked features : 98

Number of classes : 6

Settings

-----

Number of trees : 600

Max tree depth : 6

Training time (sec) : 536.9734

Training accuracy : 0.7233

Validation accuracy : 0.695

Training log\_loss : 0.5382

Validation log\_loss : 0.5738

In [17]: `error = 1 - 0.6952937280839538`

`print error`

0.304706271916

In [ ]: